

# Data Science I

Overall Points: 75 / 100

## Exercise 6: Linear Algebra & Linear Regression

Submission Deadline: January 12 2026, 07:00 UTC

University of Oldenburg

Winter 2025/2026

Instructors: Jannik Schröder, Wolfram "Wolle" Wingerath

Submitted by: <Owiti, Horata >

### Part 1: Linear Algebra

35 / 50

1.) Give a pair of square matrices A and B such that:

a)  $AB = BA$  (it commutes)

Solution:

Let

A =

$$\begin{bmatrix} 2 & 0 \\ 0 & 3 \end{bmatrix}$$

B =

$$\begin{bmatrix} 6 & 0 \\ 0 & 5 \end{bmatrix}$$

AB=

$$\begin{bmatrix} 12 & 0 \\ 0 & 15 \end{bmatrix}$$

BA=

$$\begin{bmatrix} 12 & 0 \\ 0 & 15 \end{bmatrix}$$



AB=BA therefore commutes

b)  $AB \neq BA$  (it does not commute)

**Solution:**

Let A =

$$\begin{bmatrix} 1 & 1 \\ 0 & 0 \end{bmatrix}$$

B=

$$\begin{bmatrix} 0 & 0 \\ 1 & 1 \end{bmatrix}$$

AB=

$$\begin{bmatrix} 1 & 1 \\ 0 & 0 \end{bmatrix}$$

BA=

$$\begin{bmatrix} 0 & 0 \\ 1 & 1 \end{bmatrix}$$



$AB \neq BA$  does not commute

2.) Are the matrices  $MM^T$  and  $M^T M$  square and symmetric? Please explain.

**Solution:**

Yes, the matrices are always square and symmetric.

To check if  $MM^T$  is a square, its matrix  $M = (n \cdot m)$  and  $M^T (n \cdot m)$  results into  $(n \cdot n)$  which is a square and for  $M^T M$ , its matrix  $M^T = (m \cdot n)$  ( $n \cdot m$ ) results into  $(m \cdot m)$  which is a square

To check for Symmetry For  $MM^T$ ;  $(MM^T)^T = (M^T)^T M^T = MM^T$



For  $M^T M$ ;  $(M^T M)^T = M^T (M^T)^T = M^T M$

3.) Let  $A$  be a feature matrix,  $w$  be a weight vector, and  $b$  be the target vector, so that matrix equation  $Aw = b$  holds. Please explain how to compute  $w$  with linear algebra. Does it make a difference whether  $A$  is a square and invertible matrix? Why or why not?

**Solution:**

Aw means that each row of A is multiplied with w to give one predicted value ..... X

4.) Please implement a function that computes the result of multiplying two matrices. Compare the speed of a library function for matrix multiplication to your own implementation.

a) How much faster is the library on products of random  $n \times n$  matrices, as a function of  $n$  as  $n$  gets large?

**Solution:**

```
In [2]: import numpy as np
import time

def matmul_naive(A, B):
    n, m = A.shape
    m2, p = B.shape
    assert m == m2
    C = np.zeros((n, p), dtype=float)
    for i in range(n):
        for j in range(p):
            s = 0.0
            for k in range(m):
                s += A[i, k] * B[k, j]
            C[i, j] = s
    return C

def benchmark_square(ns=(100, 200, 300, 400, 500)):
    results = []
    for n in ns:
        A = np.random.rand(n, n)
        B = np.random.rand(n, n)

        t0 = time.perf_counter()
        C1 = matmul_naive(A, B)
        t1 = time.perf_counter()

        t2 = time.perf_counter()
        C2 = A @ B
        t3 = time.perf_counter()
```

```

assert np.allclose(C1, C2, atol=1e-8)

naive_time = t1 - t0
numpy_time = t3 - t2
speedup = naive_time / numpy_time

results.append((n, naive_time, numpy_time, speedup))
return results

print(benchmark_square())

```

[(20, 0.004139457974815741, 0.0038855409948155284, 1.065349195990729),  
(40, 0.02042841599904932, 4.041698412038386e-05, 505.44137430448393), (60,  
0, 0.047660417010774836, 2.9791990527883172e-05, 1599.77283042461), (80,  
0.08621050001238473, 6.170800770632923e-05, 1397.0715182163035), (100,  
0.16643537499476224, 0.0011534999939613044, 144.2872786008402)]

**b) What about the product of an  $n \times m$  and  $m \times n$  matrix, where  $n \ll m$ ?**

**Solution:**

```

In [4]: def benchmark_rectangular(n_values=(10, 20, 30), m_multiplier=100):
    results = []
    for n in n_values:
        m = n * m_multiplier # m is much larger than n
        A = np.random.rand(n, m) # n x m matrix
        B = np.random.rand(m, n) # m x n matrix

        t0 = time.perf_counter()
        C1 = matmul_naive(A, B)
        t1 = time.perf_counter()

        t2 = time.perf_counter()
        C2 = A @ B
        t3 = time.perf_counter()

        assert np.allclose(C1, C2, atol=1e-8)

        naive_time = t1 - t0
        numpy_time = t3 - t2
        speedup = naive_time / numpy_time

        results.append((n, m, naive_time, numpy_time, speedup))
    return results

print(benchmark_rectangular())

```

[(10, 1000, 0.02686062501743436, 0.00013983299140818417, 192.09075588625  
544), (20, 2000, 0.14394954199087806, 0.0011086660088039935, 129.8403133  
5656076), (30, 3000, 0.45764037501066923, 0.00019087499822489917, 2397.5  
92032831103)]

c) Now make sure that your implementation calculates  $C = A \cdot B$  by first transposing B internally, so that all dot products are computed along rows of the matrices. Does that improve performance compared with an implementation that does not transpose B first? By how much?

**Solution:**

In [18...]

```

import numpy as np
import time

def matmul_naive(A, B):
    n, m = A.shape
    m2, p = B.shape
    assert m == m2
    C = np.zeros((n, p), dtype=float)
    for i in range(n):
        for j in range(p):
            s = 0.0
            for k in range(m):
                s += A[i, k] * B[k, j]
            C[i, j] = s
    return C

def matmul_with_transpose(A, B):
    n, m = A.shape
    m2, p = B.shape
    assert m == m2
    C = np.zeros((n, p), dtype=float)
    # Transpose B
    B_T = B.T
    for i in range(n):
        for j in range(p):
            s = 0.0
            for k in range(m):
                s += A[i, k] * B_T[j, k] # indices change with transpose
            C[i, j] = s
    return C

def benchmark_transpose_comparison(ns=(100, 200, 300)):
    results = []
    for n in ns:
        A = np.random.rand(n, n)
        B = np.random.rand(n, n)

        # Time the naive implementation
        t0 = time.perf_counter()
        C1 = matmul_naive(A, B)
        t1 = time.perf_counter()
        naive_time = t1 - t0

        # Time the implementation with transpose
        t2 = time.perf_counter()
        C2 = matmul_with_transpose(A, B)
        t3 = time.perf_counter()
        transpose_time = t3 - t2

        results.append((n, naive_time, transpose_time))
    return results

```

```

# Time NumPy's implementation for reference
t4 = time.perf_counter()
C3 = A @ B
t5 = time.perf_counter()
numpy_time = t5 - t4

# Verify all implementations give the same result
assert np.allclose(C1, C2, atol=1e-8)
assert np.allclose(C1, C3, atol=1e-8)

results.append((n, naive_time, transpose_time, numpy_time,))

return results

print(benchmark_transpose_comparison())

```

missing answer  
sentence -5 ✓

`[(100, 0.19429195899283513, 0.16745750000700355, 0.0001834580034483224  
2), (200, 1.3465008329949342, 1.3584501250006724, 0.000522125017596408  
7), (300, 4.636155791988131, 4.5665406250045635, 0.002562917012255639)]`

## Part 2: Linear Regression

40 / 50

5.) Construct an example on  $n \geq 6$  points where the optimal regression line is  $y = x$ , even though none of the input points lie directly on this line. Please visualize your example with a plot.

Solution:

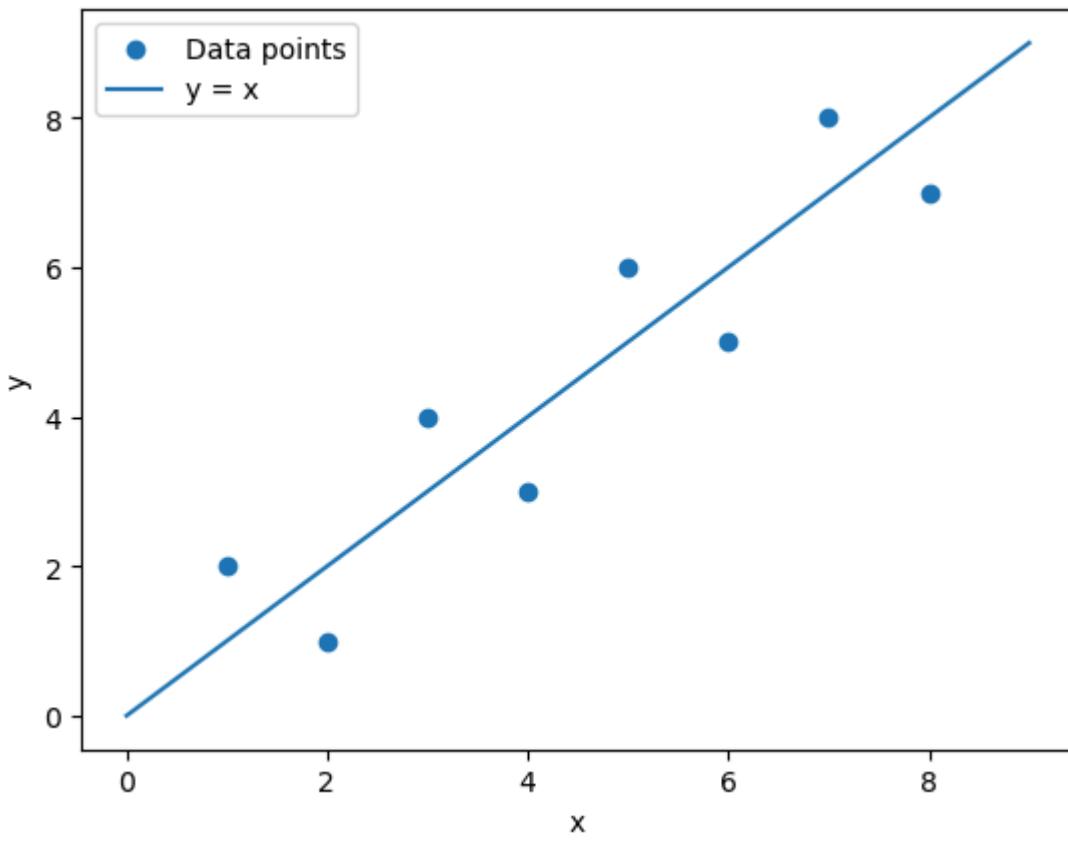
```

In [20...]: import numpy as np
          import matplotlib.pyplot as plt

          points = np.array([[1,2],[2,1],[3,4],[4,3],[5,6],[6,5],[7,8],[8,7]])
          x = points[:,0]
          y = points[:,1]

          plt.scatter(x, y, label="Data points")
          plt.plot([0,9], [0,9], label="y = x") # the regression line target
          plt.xlabel("x")
          plt.ylabel("y")
          plt.legend()
          plt.show()

```



6.) Suppose we fit a regression line to predict the shelf life of an apple based on its weight. For a particular apple, we predict the shelf life to be 4.6 days. The apples residual is -0.6 days. Did we over or under estimate the shelf-life of the apple? Please explain your reasoning.

Solution:

In [19...]

```
# Given:
predicted = 4.6          # y_hat
residual = -0.6          # r = y - y_hat

# residual = actual - predicted => actual = predicted + residual
actual = predicted + residual

print(f"Predicted shelf life (y_hat): {predicted} days")
print(f"Residual (y - y_hat): {residual} days")
print(f"Actual shelf life (y): {actual} days")

if predicted > actual:
    print("Conclusion: We OVERestimated the shelf life.")
elif predicted < actual:
    print("Conclusion: We UNDERESTIMATED the shelf life.")
else:
    print("Conclusion: Perfect prediction (no error).")
```

Predicted shelf life ( $y_{\text{hat}}$ ): 4.6 days  
 Residual ( $y - y_{\text{hat}}$ ): -0.6 days  
 Actual shelf life ( $y$ ): 3.9999999999999996 days  
 Conclusion: We OVERestimated the shelf life.



7.) Suppose we want to find the best-fitting function  $f(x) = y = w^2x + wx$ . How can we use linear regression to find the best value of  $w$ ? Please explain.

**Solution:**

In [21...]

```
import numpy as np
import matplotlib.pyplot as plt

# Example data (x, y)
x = np.array([1, 2, 3, 4, 5, 6])
y = np.array([3, 8, 15, 24, 35, 48])

# Model: y = w^2 x + w x
def predict(x, w):
    return w**2 * x + w * x

# Mean Squared Error
def mse(y, y_pred):
    return np.mean((y - y_pred)**2)

# Try many values of w
w_values = np.linspace(-5, 5, 1000)
errors = np.array([mse(y, predict(x, w)) for w in w_values])

# Best w
best_w = w_values[np.argmin(errors)]

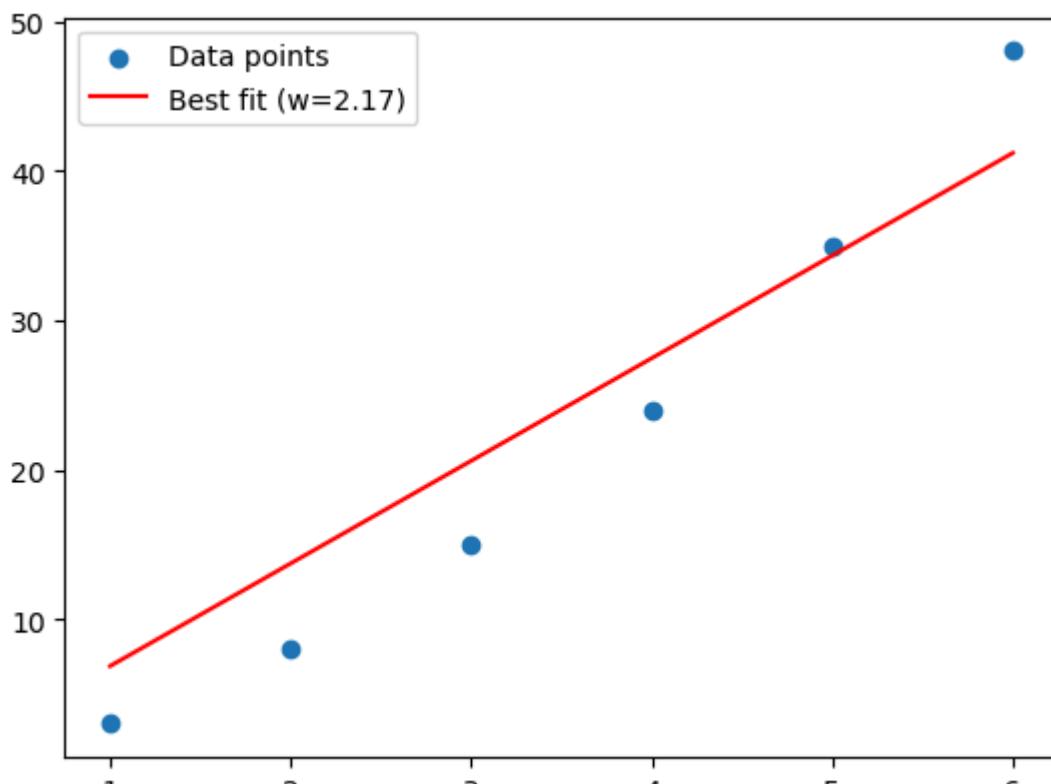
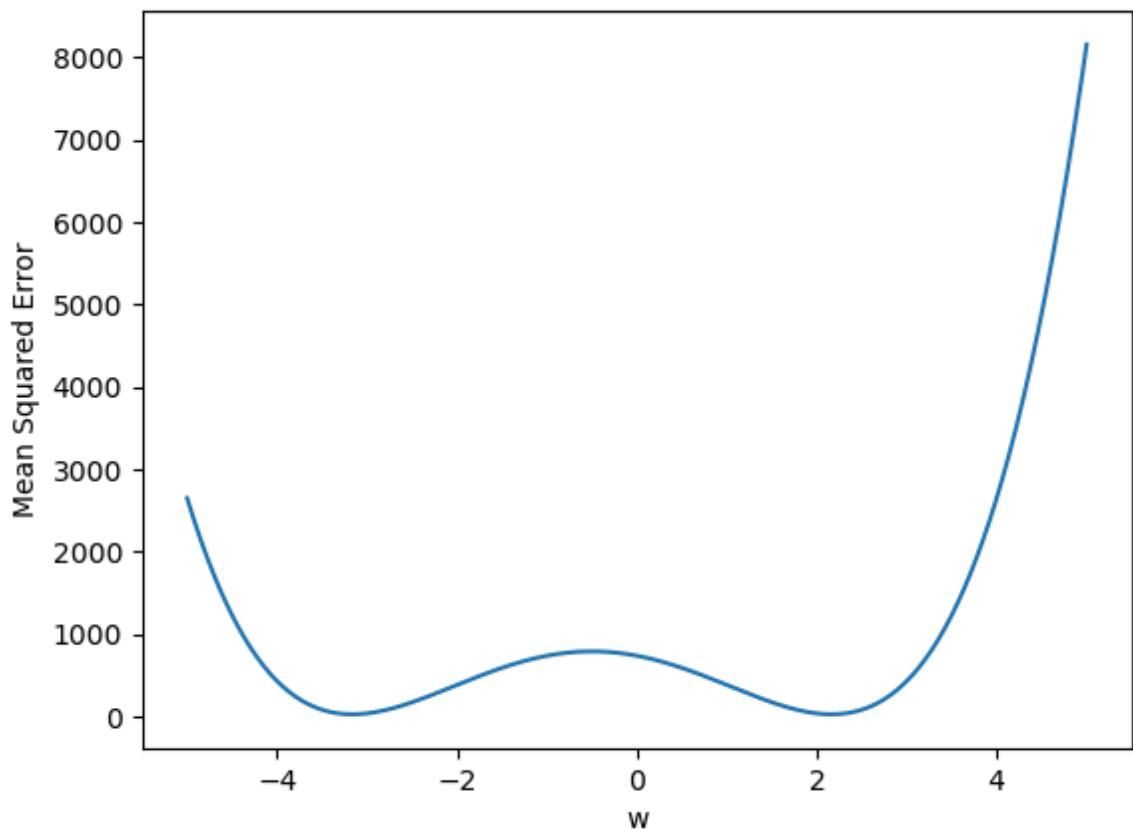
# Output
print("Best w:", best_w)
print("Minimum MSE:", np.min(errors))

# Plot error vs w
plt.plot(w_values, errors)
plt.xlabel("w")
plt.ylabel("Mean Squared Error")
plt.title("MSE as a function of w")
plt.show()

# Plot best fit
plt.scatter(x, y, label="Data points")
plt.plot(x, predict(x, best_w), color="red", label=f"Best fit (w={best_w})")
plt.legend()
plt.show()
```

Best w: 2.1671671671673  
 Minimum MSE: 22.979071347060636



**MSE as a function of w**

X

**8.) Consider the following data set:**

In [32...]:

```
import numpy as np
import matplotlib.pyplot as plt
```

```

np.random.seed(42)

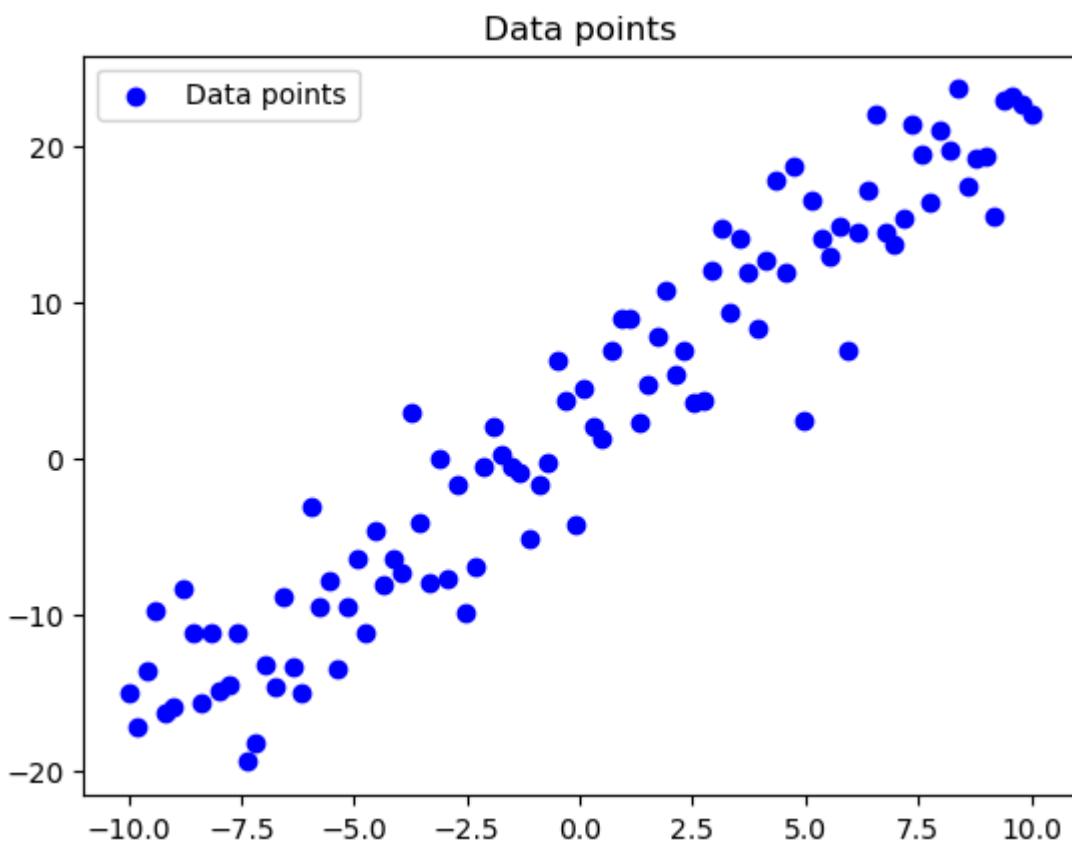
# Parameters of the line
m = 2
b = 3

# Number of data points
n = 100

# Generate data points
x = np.linspace(-10, 10, n)
y = m * x + b + np.random.normal(0, 4, n)

# Plot the data points
plt.scatter(x, y, color='blue', label='Data points')
plt.title('Data points')
plt.legend()
plt.show()

```



Assume that the data points have been generated by a linear function  $y = mx + b$  with some random noise added to the y values. Please implement the gradient descent algorithm to estimate the parameters m (slope) and b (intercept) of the linear function via the following steps.

1. Plot the provided dataset on a scatter plot.
2. Implement the gradient descent algorithm. Use the mean squared error as the cost function.
3. For each iteration of the gradient descent algorithm, plot the line with the current estimate of m and b on the scatter plot with the data points. You should end up with a plot where lines of different colors represent the

**estimated line of best fit at different stages of the algorithm. This will help to visualize how the line of best fit improves with each iteration.**

**Solution:**

In [22...]

```

import numpy as np
import matplotlib.pyplot as plt

# -----
# 1) Generate the dataset (given in the sheet)
# -----
np.random.seed(42)

m_true = 2
b_true = 3
n = 100

x = np.linspace(-10, 10, n)
y = m_true * x + b_true + np.random.normal(0, 4, n)

# -----
# 2) Gradient Descent settings
# -----
m = 0.0
b = 0.0
alpha = 0.001      # learning rate
iterations = 60     # number of steps

# Mean Squared Error
def mse(y, y_pred):
    return np.mean((y - y_pred) ** 2)

# -----
# 3) Run Gradient Descent + plot lines over time
# -----
plt.figure()
plt.scatter(x, y, label="Data points")

for t in range(iterations):
    y_pred = m * x + b

    # Gradients for MSE  $J(m,b) = (1/n) * \sum (y - (mx + b))^2$ 
    dm = (-2 / n) * np.sum(x * (y - y_pred))
    db = (-2 / n) * np.sum(y - y_pred)

    # Update parameters
    m = m - alpha * dm
    b = b - alpha * db

    # Plot the current line every few iterations (to avoid too many lines)
    if t % 5 == 0 or t == iterations - 1:
        plt.plot(x, m * x + b, label=f"iter {t}: m={m:.2f}, b={b:.2f}")

plt.title("Gradient Descent: line improves over iterations")
plt.xlabel("x")
plt.ylabel("y")

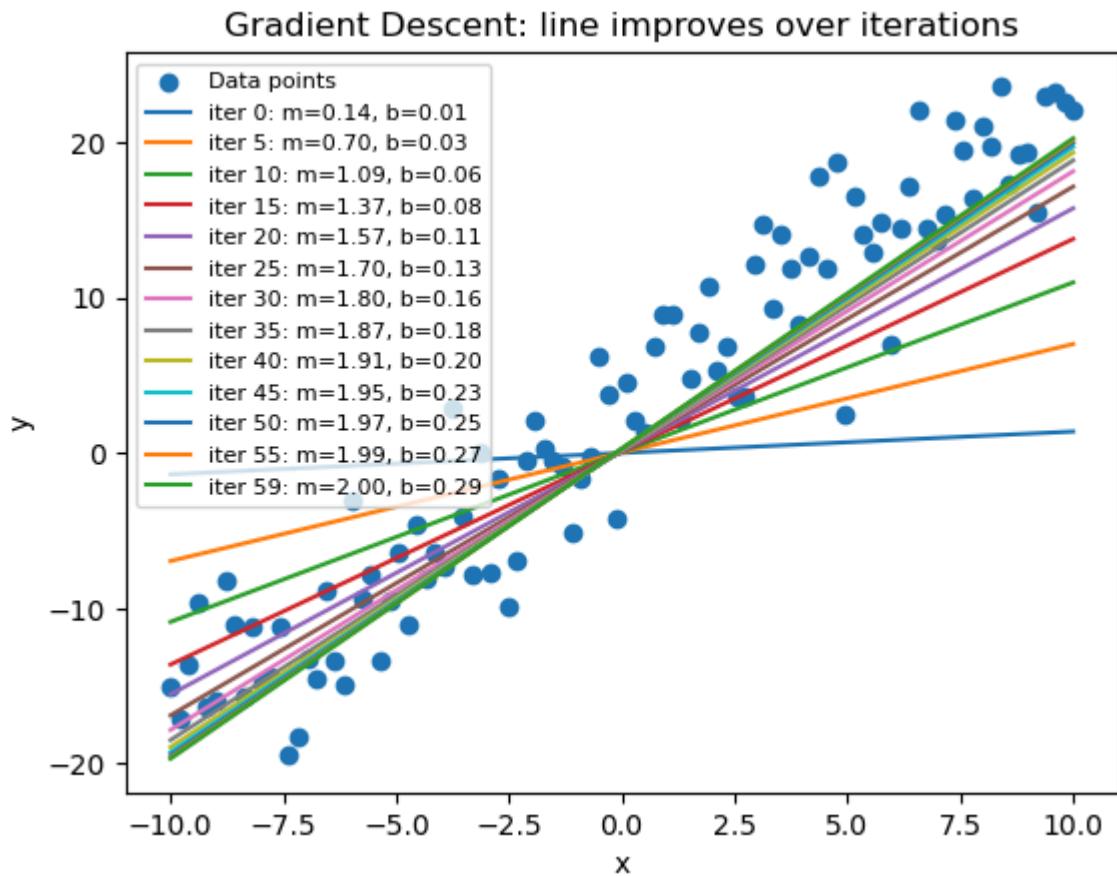
```

```

plt.legend(fontsize=8)
plt.show()

print("Final estimated m:", m)
print("Final estimated b:", b)
print("Final MSE:", mse(y, m * x + b))

```



Final estimated m: 1.997968407985234  
 Final estimated b: 0.29254244691517245  
 Final MSE: 18.321898945347083

✓

## Finally: Submission

Save your notebook and submit it (as both **notebook and PDF file**). And please don't forget to ...

- ... choose a **file name** according to convention (see Exercise Sheet 1, but please **add your group name as a suffix** like `_group01`) and to
- ... include the **execution output** in your submission!